

# Final Technical Documentation

*Salud Revenue Partners*

## Dynamic Payer Knowledge Base

### 1. Project Overview

The Salud Dynamic Payer Knowledge Base Agent is an AI-powered policy intelligence platform designed to support healthcare organizations in navigating the complexity of payer policy documentation. Built using **Streamlit** and modern **Retrieval-Augmented Generation (RAG)** architecture, the system enables users to ingest, analyze, compare, and query large volumes of payer policy PDFs (e.g., Medicaid, Medicare Advantage, and commercial payer rules).

Healthcare policy teams traditionally rely on **manual reviews, static summaries, and spreadsheet-based tracking** to interpret payer rules. This approach is time-intensive, error-prone, and poorly suited for handling frequent policy updates across multiple payers.

Salud addresses this gap by:

- Structuring unstructured policy PDFs into searchable, versioned knowledge assets
- Enabling **natural language querying** with policy-grounded answers
- Providing **page-level citations and traceability** suitable for audit and compliance contexts
- Automatically detecting and summarizing policy changes across versions

The platform is designed to support **policy analysts, compliance teams, revenue cycle managers, and healthcare consultants** who require fast, reliable, and defensible access to payer policy information.

### 2. Core Functional Capabilities

#### 2.1. Intelligent Policy Ingestion & Structuring

We implement a robust ingestion pipeline that transforms raw policy PDFs into structured, query-ready data assets.

**Key capabilities include:**

- Automated ingestion of policy PDFs organized by:
  - run\_date: Timestamp of ingestion, enabling version control

- payer\_id: Unique identifier for each payer (e.g., Medicaid state, commercial insurer)
- Page-level PDF parsing to preserve contextual granularity
- Extraction and storage of structured metadata, including:
  - Source file name
  - Page number
  - Payer identifier
  - Ingestion run date
  - Extracted text content

Each policy document is decomposed into semantically meaningful text chunks, ensuring that downstream retrieval and citation remain precise.

#### **Outcome:**

Every policy is transformed into a fully indexed, metadata-rich knowledge object, enabling fast retrieval, historical comparison, and audit-grade traceability.

## **2.2. Semantic Search & Retrieval-Augmented Generation (RAG)**

To overcome the limitations of keyword-based search, our project leverages **semantic embeddings** to capture the meaning of policy language.

#### **Technical approach:**

- Policy text is converted into dense vector embeddings using **Sentence-Transformers (MiniLM)**
- Embeddings are stored in a **Chroma vector database**, optimized for similarity search
- Queries are embedded in the same semantic space, enabling:
  - Conceptual matching (not just exact wording)
  - Retrieval of relevant policy sections even when terminology differs

#### **Advanced retrieval controls:**

- Metadata-based filtering:
  - Restrict searches to a specific payer

- Filter by ingestion date or most recent policy version
- Top-k retrieval ensures only the most relevant policy passages are passed to the LLM

**Outcome:**

Before generating any response, the system retrieves **the most contextually relevant policy excerpts**, ensuring factual grounding and minimizing hallucinations.

### 2.3. Complexity, Uncertainty, and Sense-Making in Data-Driven Projects

Our approach integrates **Anthropic Claude** as its language generation engine, operating strictly within a RAG-based constraint system.

The platform supports **two distinct response modes**:

Mode	Description	Primary Use Case
<b>Strict Mode</b>	Generates answers exclusively from retrieved policy text, without external knowledge	Legal, compliance, audit, and regulatory workflows
<b>Hybrid Mode</b>	Combines retrieved policy text with general healthcare domain knowledge (clearly labeled)	Interpretive analysis, consulting discussions

**Key safeguards:**

- All answers are grounded in retrieved policy passages
- Inline citations reference:
  - Exact PDF file
  - Page number(s)
- Responses clearly differentiate between policy-derived facts and general context (Hybrid Mode only)

**Outcome:**

Users receive **trustworthy, explainable answers** that can be defended in compliance reviews and stakeholder discussions.

### 2.4. Change Intelligence & Policy Version Comparison

Healthcare payer policies change frequently, often with subtle wording modifications that have significant downstream impact. We are incorporating a dedicated Change Intelligence Engine to address this challenge.

#### **How it works:**

- Compares the latest ingestion run against prior runs for the same payer
- Identifies:
  - Added clauses
  - Removed sections
  - Modified policy language
- Summarizes changes in clear, actionable language
- Avoids speculative interpretation by anchoring all differences in source text

#### **Use cases enabled:**

- Rapid review of policy updates
- Proactive identification of reimbursement or compliance risks
- Reduced dependency on manual redlining and document diffing

#### **Outcome:**

Policy teams can track and understand policy evolution in minutes instead of hours, improving responsiveness and decision quality.

## **2.5. Traceability, Auditing & Compliance Readiness**

Our project is designed with healthcare compliance standards in mind, emphasizing full transparency across the AI workflow.

#### **Traceability features include:**

- Direct linkage from AI-generated answer → retrieved text → original PDF
- In-app preview of cited policy pages
- Download access to original policy documents
- Display of contextual metadata (payer, run date, page number)

#### **Audit support:**

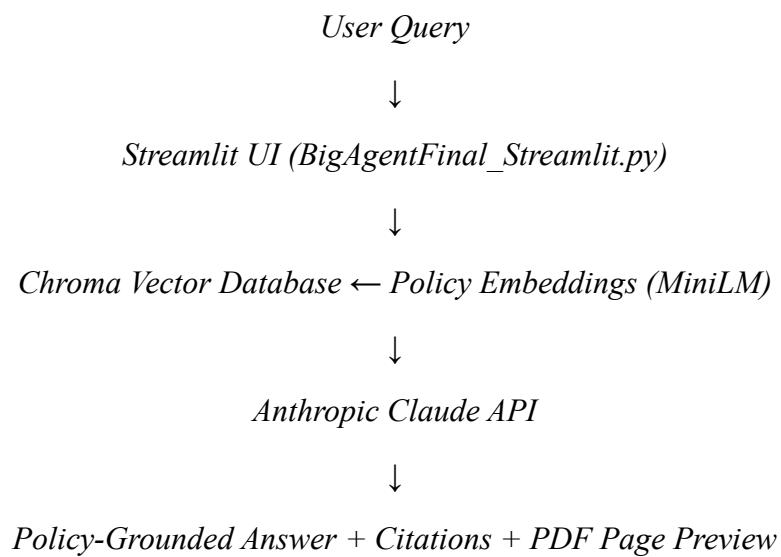
- Ensures all insights are verifiable

- Enables documentation of decision rationale
- Supports internal audits and external regulatory reviews

**Outcome:**

The system meets the expectations of **audit-ready, enterprise-grade healthcare analytics tooling.**

### 3. System Architecture Overview



#### 3.1. Core Components

Component	Role
Streamlit UI	Front-end interface for document upload, querying, and visualization
PDF Parser	Extracts page-level text and metadata from policy PDFs
MiniLM	Converts policy text into semantic embeddings
Chroma DB	Stores and retrieves embeddings for semantic search
Claude LLM	Generate grounded, context-aware responses
Change Engine	Detects and summarizes policy updates across versions

#### 4. Repository Structure & Design Rationale

File / Folder	Description
BigAgentFinal_Streamlit.py	Main application entry point managing UI, ingestion, retrieval, and response generation
requirements.txt	Python dependency list ensuring environment reproducibility
.streamlit/config.toml	UI configuration (custom Salud green theme)
Charlie Output/	Local data processing artifacts (excluded from version control)
Salud_main_1/	Root directory for all ingestion runs
Salud_main_1/<run_date>/<payer_id>/	Structured storage for versioned payer policies
README.md	Setup instructions and project overview

This structure enforces **clear separation between code, configuration, data, and outputs**, supporting maintainability and scalability.

#### 5. Local Execution Instructions

1. Clone the repository

```
git clone https://github.com/sarthakc123/BIG_KnowledgeBase.git
```

```
cd BIG_KnowledgeBase
```

## 2. Install dependencies

```
pip install -r requirements.txt
```

## 3. Launch the application

```
streamlit run BigAgentFinal_Streamlit.py
```

## 4. Configure secrets

- Navigate to **Streamlit** → **App Settings** → **Secrets**
- Add:

```
ANTHROPIC_API_KEY = "sk-ant-..."
```

## 6. Deploying Streamlit Cloud

- **Main file:** *BigAgentFinal\_Streamlit.py*
- Secrets managed through Streamlit Cloud's secure environment
- Deployed application is accessible via a shareable URL
- No sensitive credentials are stored in the repository

## 7. Data Refresh & Re-Indexing Workflow

To update or expand the policy corpus:

1. Open *navigation.ipynb*
2. Update policy URLs or document sources
3. Insert Anthropic API key (securely)
4. Execute all cells to:
  - Re-ingest PDFs
  - Recompute embeddings
  - Update vector store

This design supports **continuous policy monitoring** with minimal manual intervention.

## 8. Security, Privacy & Compliance Considerations

- No hardcoded API keys or credentials
- All secrets managed via Streamlit's encrypted secrets manager
- Processes public payer policy documents only
- No handling of PHI or PII
- Full citation trail ensures explainability and accountability

## 9. Technology Stack Summary

<b>Layer</b>	<b>Technology</b>
Interface	Streamlit
AI Engine	Anthropic Claude
Embeddings	Sentence-Transformers (MiniLM)
Vector Store	Chroma
Language	Python
Document Parsing	PyPDF, LangChain utilities
Deployment	Streamlit Community Cloud

## 10. Project Value & Differentiation

Our approach distinguishes itself by combining:

- Healthcare-domain RAG architecture
- Policy version intelligence
- Audit-grade traceability
- Practical deployment via lightweight cloud tooling

The system demonstrates how LLMs can be responsibly operationalized in regulated domains, balancing innovation with compliance, transparency, and reliability.